



Udacity Project "Explore Weather Trends"

- Author: Christian Altmooos
- First project of the "Data Analyst Nanodegree" Program of Udacity
- Summary: In this project, you will analyze local and global temperature data and compare the temperature trends where you live to overall global temperature trends.
- Tools used: SQL and jupyter notebook (python)

Get data from Udacity SQL workspace

As I'm living in Germany I will have a look on Germanies capital city **Berlin**.

- First check if "Berlin" is in the dataset

```
SQL: SELECT COUNT ( * ) FROM city_list WHERE city = 'Berlin'
```

Result: count 1

- Yes, very good. Now check number of records of city_data

```
SQL: SELECT COUNT ( * ) FROM city_data WHERE city = 'Berlin'
```

Result: count 271

- Check the number of results of global_data

```
SQL: SELECT COUNT ( * ) FROM global_data
```

Result: count 266

- So we see there is slight mismatch so it makes sense to select all city data for Berlin and global data in a full join. The data can be cleaned later if there is a need.

SQL:

```
SELECT          city.year AS year,
                city.city AS city,
                city.country AS country,
                city.avg_temp AS city_avg_temp,
                global.year AS global_year,
                global.avg_temp AS global_avg_temp
FROM            city_data city
FULL OUTER JOIN global_data global
ON              city.year = global.year
WHERE           city.city = 'Berlin'
```

Result: count 271

The result if the last SQL I downloaded to **result.csv**

Load Libraries

In [45]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

Load the data

Load the data previously downloaded into results.csv. We use pandas read_csv.

In [46]:

```
try:
    city_data = pd.read_csv('results.csv')
except Exception as er:
    print(str(er))
print(city_data.shape)
```

(271, 6)

Ok, everything all right, 271 records, 6 columns. Let's continue ...

NA handling

As we saw in the SQL session before there is mismatch of the 2 data sources. The full join might also produced NA values. So let's examine them and decide what to do.

Count NA records

In [47]:

```
city_data.isnull().sum()
```

Out[47]:

```
year          0
city          0
country       0
city_avg_temp  4
global_year    7
global_avg_temp 7
dtype: int64
```

Not too much values are missing let's have a closer look.

Display NA records

In [48]:

```
print(city_data[city_data.isnull().any(axis=1)])
```

	year	city	country	city_avg_temp	global_year	global_avg_tem
p						
0	1743	Berlin	Germany	6.33	NaN	Na
N						
1	1744	Berlin	Germany	10.36	NaN	Na
N						
2	1745	Berlin	Germany	1.43	NaN	Na
N						
3	1746	Berlin	Germany	NaN	NaN	Na
N						
4	1747	Berlin	Germany	NaN	NaN	Na
N						
5	1748	Berlin	Germany	NaN	NaN	Na
N						
6	1749	Berlin	Germany	NaN	NaN	Na
N						

Seems to be at the beginning of the time series there was an issue. The easiest way now to handle those NA records, is to get rid of them! :-)

Delete NAs from the time series

Drop those records directly in the dataset. Also drop the column `global_year` as it is so to say keeping double information.

In [49]:

```
city_data.dropna(inplace=True)
city_data.drop(labels='global_year', axis=1, inplace=True)
```

Quick check if it worked

Just to be sure count again the count of NAs.

In [50]:

```
city_data.isnull().sum()
```

Out[50]:

```
year          0
city          0
country       0
city_avg_temp 0
global_avg_temp 0
dtype: int64
```

The weather data are now adjusted and cleansed. Now we can step into the analysis.

Weather Trend Analysis

Compute the moving average

Calculate the moving average. Now the question is, how to choose the window for the moving average calculation. I was researching a bit (actually not too much). The website [climate4you.com](https://www.climate4you.com/DataSmoothing.htm) (<https://www.climate4you.com/DataSmoothing.htm>) is describing that. They have chosen a **3 years** window. However 3 years was still a bit too noisy. After some variations I decided to multiply that by 10, hence I used **30 years** as rolling window. For the calculation itself I use pandas [rolling](https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.rolling.html) (<https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.rolling.html>). The rolling method "slices" the data based on a window and one can apply any pandas method on top it. Like the mean calculation in our case.

In [51]:

```
# rolling window calculation
constWindow = 30
city_data = city_data.assign(city_mov_avg=city_data['city_avg_temp'].rolling(window=constWindow).mean())
city_data = city_data.assign(global_mov_avg=city_data['global_avg_temp'].rolling(window=constWindow).mean())
```

Summary of data

In [52]:

```
print(city_data.describe())
```

	year	city_avg_temp	global_avg_temp	city_mov_avg	\
count	264.000000	264.000000	264.000000	235.000000	
mean	1881.500000	8.917727	8.359394	8.866226	
std	76.354437	0.883601	0.575184	0.292501	
min	1750.000000	4.840000	5.780000	8.437000	
25%	1815.750000	8.357500	8.077500	8.649167	
50%	1881.500000	8.935000	8.365000	8.789333	
75%	1947.250000	9.485000	8.700000	9.017000	
max	2013.000000	10.960000	9.730000	9.851667	

	global_mov_avg
count	235.000000
mean	8.313193
std	0.348127
min	7.706333
25%	8.095667
50%	8.192333
75%	8.634500
max	9.270667

The data seems to be average temperatures of Berlin compared to the rest of the world seems to be close. But if you look closer it is different.

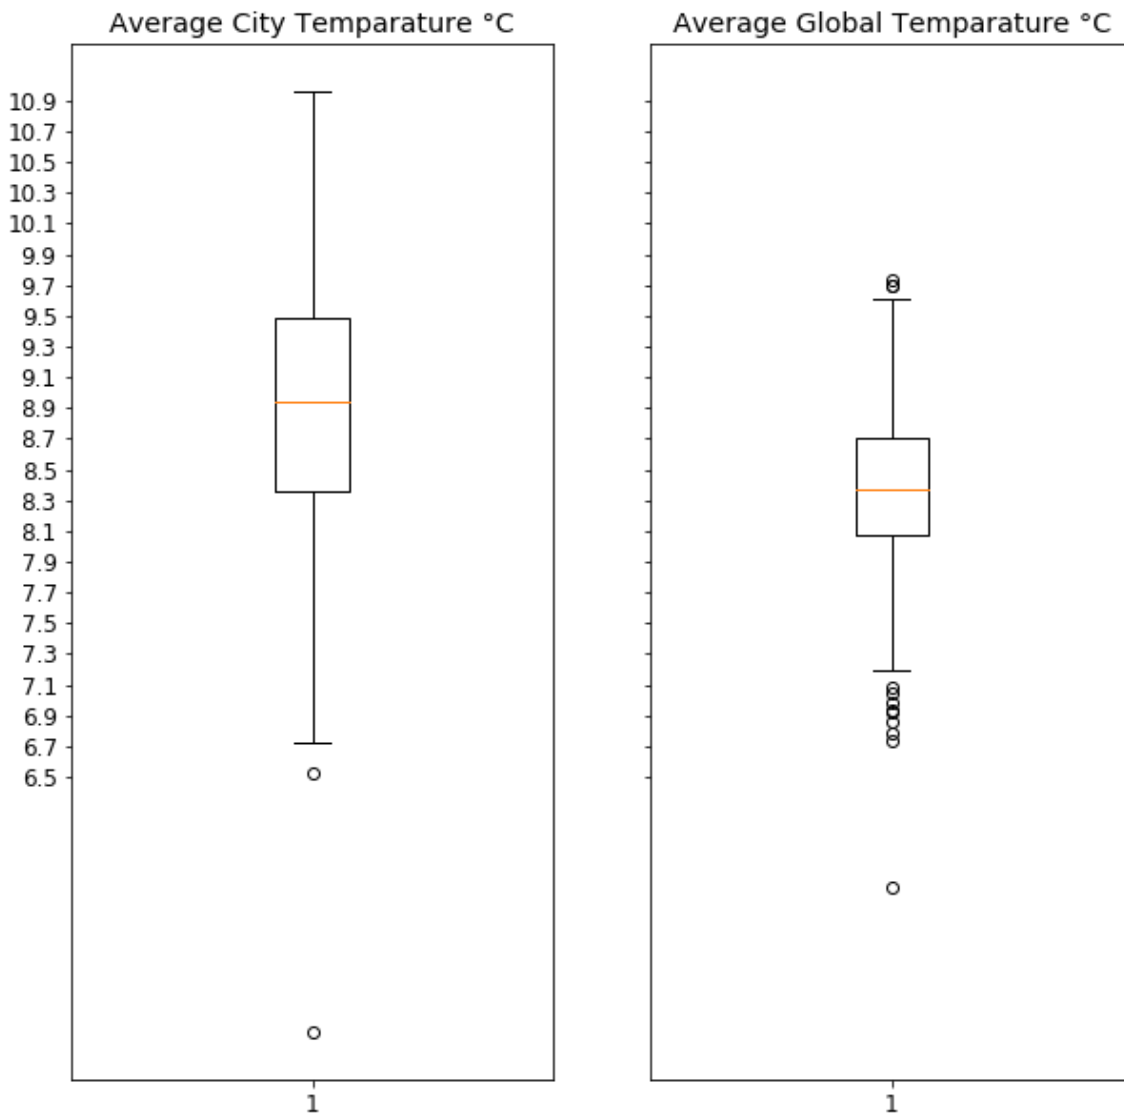
- Max temp in Berlin (city_avg_temp) 10,96°C. Compared to 9,73°C globally. More than one degree! Taking current discussions into account 1 degree increase in temperature can cause natural disasters
- The Berlin temperature (Standard Deviation: 0,85) is more variable than the Global Temperature (Standard Deviation: 0,55)

Let's have a visual look on both temperature distributions.

Plot Boxplots for Average Temperature

In [53]:

```
plt.rcParams.update({'font.size': 12})
fig, axs = plt.subplots(nrows=1, ncols=2, figsize=(10,10), sharex='col', sharey='row')
axs[0].boxplot(city_data['city_avg_temp'])
axs[0].set_title('Average City Temperature °C')
axs[0].yaxis.set_ticks(np.arange(6.5,11,0.1))
axs[1].boxplot(city_data['global_avg_temp'])
axs[1].set_title('Average Global Temperature °C')
axs[1].yaxis.set_ticks(np.arange(6.5,11,0.2))
plt.show()
```



The boxplots are underlining that fact:


- The spread of Berlin is much wider than the global one
- Globally we have far more outliers (values outside the upper and the lower whisker)
- The interquartile range (upper and lower 25% around the median) for Berlin is appr. 8,35°C - 9,45°C. The global one ranges from appr. 7,15°C - 8,7°C.

Plot the data

Let's plot the data using matplotlib.pyplot. I scatter the yearly observations and put the smoothed lines on top, to see the original values (and the spread) together with the smoothed lines.

In [54]:

```
plt.rcParams.update({'font.size': 40})
plt.figure(figsize=(45,20))
plt.plot(city_data['year'], city_data['city_mov_avg'], label = 'Berlin Moving Average Temp',linewidth=4, markersize=12)
plt.scatter(city_data['year'], city_data['city_avg_temp'], label = 'Berlin Average Temp ')
plt.plot(city_data['year'], city_data['global_mov_avg'],label = 'Global Moving Average Temp',linewidth=4, markersize=12)
plt.scatter(city_data['year'], city_data['global_avg_temp'], label = 'Global Average Temp ')
plt.legend()
plt.xlabel("Years")
plt.ylabel("Temperature (°C)")
plt.title("Berlin vs. Global Average Temperature")
plt.xticks(np.arange(min(city_data['year']), max(city_data['year']), 6))
plt.ylim(bottom = 7, top = 10.5)
plt.yticks(np.arange(7, 10.5, step=0.2))
plt.xticks(rotation=50)
plt.show()
```




Ok looks like we have an issue.

- It's clearly visible beginning of, I would say starting around 1838 there is a steady increase in temperature
- In the beginning of the 1980's the level of increase was growing dramatically
- In the early centuries approx. till the mid of 19th century there were huge variations globally.
 - e.g. the huge drop in 1838 globally to an avg. temperature of approx. 7.7 °C
 - this seems to be more stable in Berlin the drop was indeed here the opposite a slight increase
- The two time series seems to be pretty much correlated, which not necessarily mean it has a causal relationship.

Let's look to the last couple of years beginning with 1980 for the global moving average temperature.

In [55]:

```
plt.rcParams.update({'font.size': 40})
city_data = city_data.loc[(city_data['year'] >= 1980)]
temp_1980 = city_data.global_mov_avg.min()
temp_2012 = city_data.global_mov_avg.max()
temp_diff = temp_2012 - temp_1980
plt.figure(figsize=(45,20))
plt.plot(city_data['year'], city_data['global_mov_avg'], linewidth=8, markersize=12)
plt.xlabel("Years")
plt.ylabel("Temperature (°C)")
plt.title("Global Moving Average Temperature since 1980 / Increase of {} °C".format(
    str('%0.2f' % round(temp_diff, 2))))
plt.xticks(np.arange(min(city_data['year']), max(city_data['year']), 1))
plt.ylim(bottom = 8.5, top = 9.5)
plt.yticks(np.arange(8.5, 9.5, step=0.1))
plt.xticks(rotation=50)
plt.show()
```



Four Key Observations Summary

1. The city data of **Berlin** is basically **more noisy** than global data. This might be due to the fact that the global data set is kind of syndication of different temperature extremes all over the world which are balancing each other out.
2. It's a bit **warmer** in **Berlin** compared to the global average. The mean of the un-smoothed Berlin **avg. temperature** is **8,93 Celsius** compared to the Global avg. temperature of **8,36 Celsius**. The respective **maximums** are **10,96 Celsius** in Berlin and **9,73 Celsius** globally.
3. There is beginning appr. with **1838** a **steady increase** in temperature, might be the negative effects of pollution already caused by the "Industrial Revolution" (which started in the late 1700's)
4. Beginning of the **1980** the **slope** of the temperature (both locally and globally) gets **steeper**, due to us as human beings. The **increase** globally from **1980 to 2012** is **0,61 °C**.